

[android](#) / [kernel](#) / [arm64](#) / [android-9.0.0\\_r0.32](#) / [./](#) / [drivers](#) / [staging](#) / [android](#) /  
**lowmemorykiller.c**

blob: 24d2745e943778200c64996510814756854fa461 [[file](#)] [[log](#)] [[blame](#)]

```
1  /* drivers/misc/lowmemorykiller.c
2  *
3  * The lowmemorykiller driver lets user-space specify a set of memory thresholds
4  * where processes with a range of oom_score_adj values will get killed. Specify
5  * the minimum oom_score_adj values in
6  * /sys/module/lowmemorykiller/parameters/adj and the number of free pages in
7  * /sys/module/lowmemorykiller/parameters/minfree. Both files take a comma
8  * separated list of numbers in ascending order.
9  *
10 * For example, write "0,8" to /sys/module/lowmemorykiller/parameters/adj and
11 * "1024,4096" to /sys/module/lowmemorykiller/parameters/minfree to kill
12 * processes with a oom_score_adj value of 8 or higher when the free memory
13 * drops below 4096 pages and kill processes with a oom_score_adj value of 0 or
14 * higher when the free memory drops below 1024 pages.
15 *
16 * The driver considers memory used for caches to be free, but if a large
17 * percentage of the cached memory is locked this can be very inaccurate
18 * and processes may not get killed until the normal oom killer is triggered.
19 *
20 * Copyright (C) 2007-2008 Google, Inc.
21 *
22 * This software is licensed under the terms of the GNU General Public
23 * License version 2, as published by the Free Software Foundation, and
24 * may be copied, distributed, and modified under those terms.
25 *
26 * This program is distributed in the hope that it will be useful,
27 * but WITHOUT ANY WARRANTY; without even the implied warranty of
28 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
29 * GNU General Public License for more details.
30 *
31 */
32
33 #define pr_fmt(fmt) KBUILD_MODNAME ": " fmt
34
35 #include <linux/init.h>
36 #include <linux/moduleparam.h>
37 #include <linux/kernel.h>
38 #include <linux/mm.h>
39 #include <linux/oom.h>
40 #include <linux/sched.h>
41 #include <linux/swap.h>
42 #include <linux/rcupdate.h>
43 #include <linux/profile.h>
44 #include <linux/notifier.h>
```

```

45
46 static u32 lowmem_debug_level = 1;
47 static short lowmem_adj[6] = {
48     0,
49     1,
50     6,
51     12,
52 };
53
54 static int lowmem_adj_size = 4;
55 static int lowmem_minfree[6] = {
56     3 * 512,      /* 6MB */
57     2 * 1024,     /* 8MB */
58     4 * 1024,     /* 16MB */
59     16 * 1024,    /* 64MB */
60 };
61
62 static int lowmem_minfree_size = 4;
63
64 static unsigned long lowmem_deathpending_timeout;
65
66 #define lowmem_print(level, x...) \
67     do { \
68         if (lowmem_debug_level >= (level)) \
69             pr_info(x); \
70     } while (0)
71
72 static unsigned long lowmem_count(struct shrinker *s,
73                                  struct shrink_control *sc)
74 {
75     return global_page_state(NR_ACTIVE_ANON) +
76           global_page_state(NR_ACTIVE_FILE) +
77           global_page_state(NR_INACTIVE_ANON) +
78           global_page_state(NR_INACTIVE_FILE);
79 }
80
81 static unsigned long lowmem_scan(struct shrinker *s, struct shrink_control *sc)
82 {
83     struct task_struct *tsk;
84     struct task_struct *selected = NULL;
85     unsigned long rem = 0;
86     int tasksize;
87     int i;
88     short min_score_adj = OOM_SCORE_ADJ_MAX + 1;
89     int minfree = 0;
90     int selected_tasksize = 0;
91     short selected_oom_score_adj;
92     int array_size = ARRAY_SIZE(lowmem_adj);
93     int other_free = global_page_state(NR_FREE_PAGES) - totalreserve_pages;
94     int other_file = global_page_state(NR_FILE_PAGES) -
95                     global_page_state(NR_SHMEM) -
96                     total_swapcache_pages();
97

```

```

98     if (lowmem_adj_size < array_size)
99         array_size = lowmem_adj_size;
100    if (lowmem_minfree_size < array_size)
101        array_size = lowmem_minfree_size;
102    for (i = 0; i < array_size; i++) {
103        minfree = lowmem_minfree[i];
104        if (other_free < minfree && other_file < minfree) {
105            min_score_adj = lowmem_adj[i];
106            break;
107        }
108    }
109
110    lowmem_print(3, "lowmem_scan %lu, %x, ofree %d %d, ma %hd\n",
111                sc->nr_to_scan, sc->gfp_mask, other_free,
112                other_file, min_score_adj);
113
114    if (min_score_adj == OOM_SCORE_ADJ_MAX + 1) {
115        lowmem_print(5, "lowmem_scan %lu, %x, return 0\n",
116                    sc->nr_to_scan, sc->gfp_mask);
117        return 0;
118    }
119
120    selected_oom_score_adj = min_score_adj;
121
122    rcu_read_lock();
123    for_each_process(tsk) {
124        struct task_struct *p;
125        short oom_score_adj;
126
127        if (tsk->flags & PF_KTHREAD)
128            continue;
129
130        p = find_lock_task_mm(tsk);
131        if (!p)
132            continue;
133
134        if (task_lmk_waiting(p) &&
135            time_before_eq(jiffies, lowmem_deathpending_timeout)) {
136            task_unlock(p);
137            rcu_read_unlock();
138            return 0;
139        }
140        oom_score_adj = p->signal->oom_score_adj;
141        if (oom_score_adj < min_score_adj) {
142            task_unlock(p);
143            continue;
144        }
145        tasksize = get_mm_rss(p->mm);
146        task_unlock(p);
147        if (tasksize <= 0)
148            continue;
149        if (selected) {
150            if (oom_score_adj < selected_oom_score_adj)

```

```

151         continue;
152         if (oom_score_adj == selected_oom_score_adj &&
153             tasksize <= selected_tasksize)
154             continue;
155     }
156     selected = p;
157     selected_tasksize = tasksize;
158     selected_oom_score_adj = oom_score_adj;
159     lowmem_print(2, "select '%s' (%d), adj %hd, size %d, to kill\n",
160                 p->comm, p->pid, oom_score_adj, tasksize);
161 }
162 if (selected) {
163     task_lock(selected);
164     send_sig(SIGKILL, selected, 0);
165     if (selected->mm)
166         task_set_lmk_waiting(selected);
167     task_unlock(selected);
168     lowmem_print(1, "Killing '%s' (%d), adj %hd,\n"
169                 "    to free %ldkB on behalf of '%s' (%d) because\n"
170                 "    cache %ldkB is below limit %ldkB for oom_score_adj %hd\n"
171                 "    Free memory is %ldkB above reserved\n",
172                 selected->comm, selected->pid,
173                 selected_oom_score_adj,
174                 selected_tasksize * (long)(PAGE_SIZE / 1024),
175                 current->comm, current->pid,
176                 other_file * (long)(PAGE_SIZE / 1024),
177                 minfree * (long)(PAGE_SIZE / 1024),
178                 min_score_adj,
179                 other_free * (long)(PAGE_SIZE / 1024));
180     lowmem_deathpending_timeout = jiffies + HZ;
181     rem += selected_tasksize;
182 }
183
184 lowmem_print(4, "lowmem_scan %lu, %x, return %lu\n",
185             sc->nr_to_scan, sc->gfp_mask, rem);
186 rcu_read_unlock();
187 return rem;
188 }
189
190 static struct shrinker lowmem_shrinker = {
191     .scan_objects = lowmem_scan,
192     .count_objects = lowmem_count,
193     .seeks = DEFAULT_SEEKS * 16
194 };
195
196 static int __init lowmem_init(void)
197 {
198     register_shrinker(&lowmem_shrinker);
199     return 0;
200 }
201 device_initcall(lowmem_init);
202
203 /*

```

```
204  * not really modular, but the easiest way to keep compat with existing
205  * bootargs behaviour is to continue using module_param here.
206  */
207 module_param_named(cost, lowmem_shrinker.seek, int, S_IRUGO | S_IWUSR);
208 module_param_array_named(adj, lowmem_adj, short, &lowmem_adj_size,
209                          S_IRUGO | S_IWUSR);
210 module_param_array_named(minfree, lowmem_minfree, uint, &lowmem_minfree_size,
211                          S_IRUGO | S_IWUSR);
212 module_param_named(debug_level, lowmem_debug_level, uint, S_IRUGO | S_IWUSR);
213
```